Machine learning: prediction, classification and clustering

UBB Faculty of Sociology

Course Agenda

#1 Intro, Simple Linear Regression

#2 Python recap, Git, Handling data, EDA

#3 Regression, Decision Trees

#4 Bias Variance, Overfitting, Classification, Metrics

#5 Random Forest Classifier, Clustering

#6 Neural Networks

#7 Help Final Project

#8 Help Final Project

2. Linear Regression, Decision Trees

#1.1 Linear Regression recap

#1.2 Decision Trees Regression

#1.3 Regression implementation

#1.4 Final Project Task 2 - Census Data EDA

#1.5 Questions & Further reading

1.1 Linear Regression recap

Predict calories

Eq of the line:
y = ax + b

Where:
a: slope
b: intercept
x: data

How lines work:
Calories =
Slope * Weight +
Intercept

Predict the calories

Calories for 15 apple pie slices:

148, 154, 167, 176, 176,

191, 203, **207**, 227, 230,

245, 251, 259, 290, 503

Mean: 228,47

Median: 207

Predict calories

Weight vs. Calories with Prediction and Errors
RMSE = 48.04, MSE = 2308.14

Legend:
- Data Points
- Regression Line
- Mean = 228.47
- Median = 207.00
- x Prediction for 20g = 129.96

Eq of the line:
y = ax + b

Where:
a: slope
b: intercept
x: data

How lines work:
Calories =
1.19*20 + 106.2=131

Code on Colab

Data is not always linear

No problem!

There are non-linear regression models:

- Decision Tree Regression

- SVM with non-linear kernels

- Neural networks

- KNN

- Ensemble methods

1.2 Decision Trees Regression

Decision Tree Regression

Decision Trees Regression is a very fast algorithm How it works: A two step repeating process

1. Take some random samples in the subtree split the dataset in two
2. Save the random split which separates the data by value the best

Then, repeat for the newly created subtrees

Decision Tree Regression

Decision Trees Regression is a very fast algorithm How it works: A two step repeating process

1. Step 1
    1. Take some random samples in the subtree split the dataset in two
    2. Save the random split which separates the data by value the best

2. Step 2
    1. Take some random samples in the subtree split the dataset in two
    2. Save the random split which separates the data by value the best

Decision Tree Regression

Decision Trees Regression is a very fast algorithm How it works: A two step repeating process

3. Step 3
    1. Take some random samples in the subtree split the dataset in two
   2. Save the random split which separates the data by value the best

         ……..

n. Step n

    n.1 Take some random samples in the subtree split the dataset in two

    n.2 Save the random split which separates the data by value the best

Decision Tree Regression

Stop Criteria:

- **Maximum Depth:** Stops growing the tree when the maximum depth is reached.

- **Minimum Samples per Leaf:** Requires a minimum number of samples in each leaf node.

- **Minimum Samples per Split:** Requires a minimum number of samples to attempt a split.

- **Maximum Number of Nodes:** Limits the total number of nodes in the tree.

- **No Further Split Improves Performance:** Stops when no split improves the target metric.

- **Early Stopping with Cross-Validation:** Stops if validation performance does not improve.

- **Insufficient Data:** Stops if a node has too few samples for further splitting.

Decision Tree Regression

Advantages:

- We get a nice, explainable model, easy to understand by humans
- Very fast prediction for new samples

Decision Tree Regression

Disadvantages:

- Leaves grow exponentially by depth
- Very prone to overfitting

Decision Tree Regression Calories

Decision Tree Regression

Code on Colab

Decision Tree Regression Calories

Decision Tree for Regression

[Code on Colab](#)

Weight <= 190.0
squared_error = 6979.316
samples = 15
value = 228.467

True   False

Weight <= 110.0
squared_error = 1709.837
samples = 14
value = 208.857

squared_error = 0.0
samples = 1
value = 503.0

Weight <= 68.5
squared_error = 404.938
samples = 8
value = 177.75

Weight <= 175.0
squared_error = 439.222
samples = 6
value = 250.333

squared_error = 130.56
samples = 5
value = 164.2

squared_error = 46.222
samples = 3
value = 200.333

squared_error = 149.44
samples = 5
value = 242.4

squared_error = 0.0
samples = 1
value = 290.0

1.3 Regression implementation

- Penguins size

- California Housing

1.4 Final Project Task 2 - Census Data EDA

https://github.com/zahariesergiu/ubb-sociology-ml/blob/main/final_project/Final_Project_Task_2_Census_Data_EDA.ipynb

1.5 Further Reading & Questions

#1 Regression Trees explained: https://www.youtube.com/watch?v=g9c66TUylZ4

#2 Decision Tees Split Criteria: https://scientistcafe.com/ids/splitting-criteria

#3 Self-supervised learning: https://ai.meta.com/blog/self-supervised-learning-the-dark-matter-of-intelligence/

#4 Feature scaling @ Kaggle: https://www.kaggle.com/code/mysarahmadbhat/all-about-feature-scaling/notebook

#5 Nature article: statistics and machine learning:
https://www.nature.com/articles/nmeth.4642?source=post_page-----64b49f07ea3-----------------------------

#6 If correlation doesn't imply causation, then what does?:
https://michaelnielsen.org/ddi/if-correlation-doesnt-imply-causation-then-what-does/

#7 Requirements for talking about cause and effect:
https://towardsdatascience.com/are-you-guilty-of-using-the-word-experiment-incorrectly-9068baeab7a4

Thank you !!

Machine Learning Engineer / Data Scientist
zahariesergiu@gmail.com
https://www.linkedin.com/in/zahariesergiu/
https://github.com/zahariesergiu/ubb-sociology-ml